

Securing Neural Interfaces: Architecture, Threat Taxonomy, and Neural Impact Scoring for Brain-Computer Interfaces

Kevin L. Qi

Qinnovate

kevin@qinnovate.com

February 2026

Abstract

Brain-computer interfaces (BCIs) are transitioning from experimental neuroscience tools to commercially deployed medical devices, with companies including Neuralink, Synchron, Blackrock Neurotech, and Paradromics advancing toward regulatory approval and new entrants such as Merge Labs raising \$252M in seed funding [2]. Yet no security framework exists that accounts for the unique risks of devices that read and write neural signals. The Common Vulnerability Scoring System (CVSS v4.0), the industry standard for vulnerability assessment, cannot express biological tissue damage, cognitive integrity violations, consent boundaries, damage reversibility, or neuroplastic consequences—dimensions critical to neural device security.

We present an integrated security framework comprising four contributions: (1) an **11-band hourglass architecture** mapping attack surfaces from neocortex to wireless radio across neural, interface, and synthetic zones; (2) **TARA**, a threat taxonomy of 102 techniques across 15 tactics and 8 domains, each classified by status, severity, and dual-use therapeutic potential; (3) **NISS**, the Neural Impact Scoring System—a CVSS v4.0 extension adding five neural-specific metrics (Biological Impact, Cognitive Integrity, Consent Violation, Reversibility, Neuroplasticity) designed to conform with FIRST.org’s official extension mechanism; and (4) the **Neural Impact Chain**, a methodology mapping security vulnerabilities to DSM-5-TR psychiatric diagnoses through a six-stage pipeline.

Analysis of all 102 techniques reveals that 96.1% require NISS extension metrics that CVSS cannot express. The Neural Impact Chain maps all techniques to 15 unique DSM-5-TR diagnostic codes across 5 psychiatric clusters, with 51 techniques posing direct diagnostic risk. The framework identifies 77 techniques (75.5%) with confirmed or probable therapeutic analogs, establishing a dual-use atlas where every attack mechanism that can harm neural tissue has a corresponding clinical application. The complete framework, threat registry, and scoring system are released as open source under the Apache 2.0 license.

Keywords: brain-computer interface, neurosecurity, CVSS, threat taxonomy, neural scoring, neuroethics, DSM-5-TR, dual-use

Contents

1	Introduction	4
1.1	The Scoring Gap	4
1.2	Contributions	4
2	Related Work	5
2.1	Foundational Neurosecurity	5
2.2	Empirical Attack Demonstrations	5

2.3	Emerging Frameworks	5
2.4	Neuroethics and Policy	6
2.5	How QIF Differs	6
3	The Hourglass Architecture	6
3.1	Design Rationale	6
3.2	Neural Band Derivation	7
3.3	Band Definitions	7
3.4	Zone Structure	8
3.5	Security Properties	8
4	TARA: Threat Taxonomy	9
4.1	Taxonomy Structure	10
4.2	Evidence Classification	10
4.3	Severity Distribution	10
4.4	Category Breakdown	11
4.5	Dual-Use Mapping	11
4.6	Representative Techniques	11
5	NISS: Neural Impact Scoring System	12
5.1	Gap Analysis: Why CVSS Alone Is Insufficient	12
5.2	Extension Metrics	13
5.2.1	BI: Biological Impact	13
5.2.2	CG: Cognitive Integrity	13
5.2.3	CV: Consent Violation	13
5.2.4	RV: Reversibility	14
5.2.5	NP: Neuroplasticity	14
5.3	PINS Flag	14
5.4	Scoring Formula	14
5.5	Vector Format	15
5.6	NISS Severity Distribution	15
6	Neural Impact Chain	16
6.1	Pipeline Architecture	16
6.2	NISS-to-DSM Bridge	17
6.3	Coverage Statistics	17
6.4	Diagnostic Cluster Distribution	17
6.5	Risk Classification	18
6.6	Example Walkthrough	18
7	Governance and Neuroethics	19
7.1	Consent Tiers	19
7.2	UNESCO Alignment	19
7.3	Neurorights Framework Integration	20
7.4	Regulatory Mapping	20
8	Case Studies	21
8.1	Scoring Comparison	21
8.2	Case 1: Cortical Signal Injection (QIF-T0001)	21
8.3	Case 2: P300 Side-Channel (QIF-T0015)	21
8.4	Case 3: Calibration Poisoning (QIF-T0034)	22
8.5	Case 4: Covert Neural Surveillance (QIF-T0050)	22

8.6 Case 5: Real-World Vulnerability Disclosure	22
9 Limitations and Future Work	23
9.1 No Empirical Validation on Real BCI Devices	23
9.2 DSM-5-TR Mapping Not Clinically Validated	23
9.3 NISS Weights Not Calibrated	23
9.4 No Interrater Reliability Study	23
9.5 Taxonomy Completeness	23
9.6 Single-Author Bias	23
9.7 AI Tool Disclosure	24
9.8 Future Work	24
10 Conclusion	24

1 Introduction

Brain-computer interfaces (BCIs) are no longer confined to research laboratories. Neuralink has implanted its N1 device in human patients [22], Synchron’s Stentrode has demonstrated motor neuroprostheses via neurointerventional surgery [25], and Blackrock Neurotech’s Utah array has enabled high-performance speech neuroprostheses [34]. Paradromics is advancing toward high-bandwidth cortical interfaces for paralyzed patients. Investment in neurotechnology grew 700% between 2014 and 2021, with the global market projected to reach \$25 billion by 2030 [30].

Despite this rapid commercialization, no security standard exists specifically for neural devices. The FDA’s 2023 cybersecurity guidance for medical devices [32] and Section 524B of the FD&C Act [33] address connected medical devices generally, but contain no provisions for the unique risks of devices that read and write neural signals. The EU Medical Device Regulation [8] similarly lacks neural-specific security requirements.

1.1 The Scoring Gap

The Common Vulnerability Scoring System (CVSS v4.0) [9] is the industry standard for assessing vulnerability severity. CVSS excels at capturing exploitability characteristics—attack vector, complexity, privileges required, user interaction—and information system impact across confidentiality, integrity, and availability. However, CVSS was designed for information technology assets. It has no concept of:

- **Biological tissue damage**—seizure induction, neural tissue necrosis, involuntary motor activation
- **Cognitive integrity**—thought privacy, perception manipulation, identity modification
- **Consent boundaries**—the distinction between operating within consented parameters and covert neural manipulation
- **Damage reversibility**—IT assets can be restored from backup; neural tissue cannot be rebooted
- **Neuroplastic consequences**—prolonged adversarial stimulation can cause lasting structural changes to the brain

When a vulnerability in a BCI device can induce seizures, decode private thoughts, or cause irreversible brain damage, a CVSS base score alone is insufficient. The gap is not theoretical: our analysis of 102 BCI attack techniques shows that 96.1% require scoring dimensions that CVSS cannot express.

1.2 Contributions

This paper presents an integrated security framework with four contributions:

1. **The Hourglass Architecture** (Section 3): An 11-band model mapping attack surfaces across three zones—neural (7 bands, from neocortex to spinal cord), interface (1 band, the electrode-tissue boundary), and synthetic (3 bands, from analog electronics to wireless radio). The hourglass shape reflects a natural security chokepoint at the neural interface.
2. **TARA Threat Taxonomy** (Section 4): A registry of 102 techniques across 15 tactics and 8 domains, each classified by evidence status, severity, and dual-use therapeutic potential. TARA is an independent taxonomy inspired by the structural methodology of MITRE ATT&CK®, tailored to the BCI domain.
3. **NISS Scoring System** (Section 5): The Neural Impact Scoring System—a CVSS v4.0 extension designed to conform with FIRST.org’s official extension mechanism (User Guide §3.11) [10], adding five neural-specific metrics. Formal registration with FIRST.org is planned as future work.
4. **The Neural Impact Chain** (Section 6): A methodology mapping security vulnerabilities to psychiatric diagnoses through a six-stage pipeline: technique → hourglass band → neural

structure → cognitive function → NISS score → DSM-5-TR diagnostic code [1].

Sections 7–8 address governance alignment with international neuroethics frameworks and present comparative case studies. Section 9 discusses limitations and validation gaps honestly. The complete framework, threat registry, and scoring system are released as open source under the Apache 2.0 license.

2 Related Work

Research at the intersection of cybersecurity and neurotechnology has progressed through three phases: foundational framing, empirical demonstration, and emerging policy response. QIF builds on each while addressing gaps left by prior work.

2.1 Foundational Neurosecurity

Denning, Matsuoka, and Kohno [6] coined the term “neurosecurity” and established the field by analyzing attack surfaces in implantable neurostimulators. Their work identified wireless reprogramming, battery depletion, and signal injection as threat categories, but predated modern high-bandwidth BCIs and did not propose a scoring system.

Bonaci, Calo, and Chizeck [3] extended this with “App Stores for the Brain,” examining privacy threats from third-party BCI applications. They identified the need for access control at the neural data layer—a concern our governance framework addresses through consent tiers (Section 7).

2.2 Empirical Attack Demonstrations

Martinovic et al. [19] demonstrated at USENIX Security 2012 that consumer-grade EEG headsets could be exploited via P300 event-related potentials to extract private information—PIN numbers, bank details, and personal preferences—through subliminal visual stimuli embedded in a game. This remains the most influential empirical demonstration of a BCI side-channel attack. Our TARA taxonomy (Section 4) classifies this as a confirmed technique in the Signal Eavesdropping category.

Ienca and Haselager [13] broadened the scope to “hacking the brain,” analyzing BCI security through the lens of informational and physical integrity. They proposed that BCI security requires both traditional cybersecurity measures and novel neuroethical safeguards—a dual requirement that QIF implements through the parallel CVSS/NISS scoring architecture.

Meng et al. [20] established an adversarial robustness benchmark for EEG-based BCIs, evaluating multiple adversarial defense approaches across neural network architectures and EEG datasets. Their work focuses on machine learning adversarial robustness; TARA extends coverage to physical-layer, protocol-layer, and cognitive-layer threats.

2.3 Emerging Frameworks

Schroder et al. [28] published the most recent analysis (2025), developing a threat model for network-based attacks on next-generation BCIs and identifying associated regulatory changes needed to address them. Their work contributes threat categorization across network pathways but does not propose a scoring system, formal threat taxonomy with technique-level granularity, or clinical impact mapping.

Camara et al. [4] and Rushanan et al. [27] surveyed security and privacy in implantable medical devices more broadly, covering pacemakers, insulin pumps, and cochlear implants alongside neurostimulators. Halperin et al. [11] demonstrated practical wireless attacks against pacemakers, establishing that implanted medical device security is not theoretical. These works

address the broader medical device landscape; QIF focuses specifically on the unique challenges of bidirectional neural interfaces.

2.4 Neuroethics and Policy

Ienca and Andorno [12] proposed four fundamental neurorights: cognitive liberty, mental privacy, mental integrity, and psychological continuity. Yuste et al. [35] added equal access and protection from algorithmic bias. These philosophical frameworks establish *what* must be protected; QIF provides the technical architecture for *how*.

Muñoz et al. [23] and Lázaro-Muñoz et al. [18] contributed empirical grounding through researcher interviews, documenting clinician concerns about adaptive DBS ethics (72% cited uncertainty about risks) and post-trial access obligations for implanted neural devices.

UNESCO’s 2025 Recommendation on the Ethics of Neurotechnology [30]—adopted by 194 Member States—represents the first global normative framework for neurotechnology governance. The OECD had previously published a Recommendation on Responsible Innovation in Neurotechnology [24] covering 36 member countries. Chile became the first nation to constitutionally protect neurorights [26]. These policy instruments establish principles and values; QIF demonstrates technical implementation (Section 7).

2.5 How QIF Differs

Prior work has addressed BCI security, neuroethics, and vulnerability scoring separately. QIF is the first framework to integrate all four layers into a single system:

1. **Architecture:** A formal band model (vs. informal threat lists)
2. **Taxonomy:** Technique-level granularity with evidence classification (vs. category-level surveys)
3. **Scoring:** A CVSS-compatible extension with neural-specific metrics (vs. qualitative risk ratings)
4. **Clinical mapping:** A pipeline from security vulnerability to psychiatric diagnosis (novel; no prior work)

The Neural Impact Chain (Section 6) has no precedent in either cybersecurity or neuroethics literature. To our knowledge, no prior work has systematically mapped vulnerability severity to DSM-5-TR diagnostic codes.

3 The Hourglass Architecture

The QIF architecture maps the full BCI attack surface using an 11-band hourglass model organized into three zones. The model is derived from two independent design traditions: the OSI networking reference model [36], which stratifies communication systems into functional layers, and functional neuroanatomy, which organizes the nervous system by anatomical structure and physiological role. The hourglass shape borrows from the Internet protocol architecture [5], where IP serves as a narrow waist through which all traffic must pass.

3.1 Design Rationale

A BCI system spans from high-level cortical computation to low-level radio transmission. Any security framework must account for threats at every point along this path. We observe that the neural interface—the electrode-tissue boundary—forms a natural chokepoint analogous to IP in the Internet hourglass: all signals, whether neural or synthetic, must cross this boundary. Above the interface, attack surfaces expand through the complexity of neural tissue. Below it, they expand through electronic and wireless systems. The resulting shape is an asymmetric hourglass with the narrowest point at the interface.

3.2 Neural Band Derivation

The seven neural bands (N1–N7) are not arbitrary divisions. They correspond directly to the canonical central nervous system hierarchy established in foundational neuroscience texts [15]: spinal cord, brainstem, cerebellum, diencephalon, basal ganglia, limbic system, and neocortex. This mapping was chosen for three reasons:

- Functional differentiation:** Each CNS division performs computationally distinct roles. The spinal cord (N1) mediates reflexes; the brainstem (N2) controls vital autonomic functions; the neocortex (N7) performs executive cognition and language. Attacks targeting different divisions produce qualitatively different harms—from respiratory arrest (N2) to identity manipulation (N7).
- Determinacy gradient:** Ascending the neural hierarchy, system behavior becomes progressively less deterministic. Spinal reflexes (N1) are largely predictable; basal ganglia dynamics (N5) exhibit chaotic sensitivity; prefrontal deliberation (N7) involves the highest degree of unpredictability. This gradient has direct security implications: attacks at lower bands produce more predictable (and often more immediately dangerous) effects, while attacks at higher bands produce less predictable but potentially more insidious cognitive consequences.
- Device deployment mapping:** Existing BCI devices map cleanly onto specific bands. Sacral nerve stimulators and spinal cord stimulators operate at N1, deep brain stimulation systems (e.g., Medtronic Percept) target the subthalamic nucleus at N5, and cortical implants (e.g., Neuralink N1 chip, Blackrock Utah array) interface at N7. The band structure ensures that every deployed device class has a defined security context.

Each band also exhibits characteristic neural oscillation frequency ranges (delta through gamma), though the bands are defined by neuroanatomy rather than frequency. The oscillation profile of each band informs signal-plausibility checks: a stimulation signal injected at N4 (thalamus) can be validated against the expected spectral envelope for thalamocortical rhythms, providing a physics-based anomaly detection mechanism.

3.3 Band Definitions

Table 1 defines all 11 bands. The model uses a 7-1-3 asymmetric structure: seven neural bands (N7–N1), one interface band (I0), and three synthetic bands (S1–S3).

Table 1: QIF Hourglass: 11 bands across three zones.

Band	Name	Zone	Attack Surface
N7	Neocortex	Neural	PFC, M1, V1, Broca, Wernicke—executive function, language, movement, perception
N6	Limbic System	Neural	Hippocampus, amygdala, insula—emotion, memory, interoception
N5	Basal Ganglia	Neural	Striatum, STN, substantia nigra—motor selection, reward, habit
N4	Diencephalon	Neural	Thalamus, hypothalamus—sensory gating, consciousness relay
N3	Cerebellum	Neural	Cerebellar cortex, deep nuclei—motor coordination, timing
N2	Brainstem	Neural	Medulla, pons, midbrain—vital functions, arousal, reflexes
N1	Spinal Cord	Neural	Cervical through sacral—reflexes, peripheral relay
I0	Neural Interface	Interface	Electrode-tissue boundary—the hourglass waist; all signals must cross
S1	Analog / Near-Field	Synthetic	Amplification, ADC, near-field EM (0–10 kHz)
S2	Digital / Telemetry	Synthetic	Decoding, BLE/WiFi, telemetry (10 kHz–1 GHz)
S3	Radio / Wireless	Synthetic	RF, directed energy, application layer (1 GHz+)

3.4 Zone Structure

Neural Zone (N7–N1). The neural zone encompasses all biological structures from the neocortex to the spinal cord. Attack surfaces in this zone involve direct interaction with neural tissue: signal injection, neural manipulation, cognitive exploitation, and physical safety threats. The ordering follows neuroanatomical hierarchy—higher bands represent more complex cognitive functions, while lower bands represent more fundamental physiological processes. Attacks at lower neural bands (N1–N2) threaten vital functions; attacks at higher bands (N6–N7) threaten cognition, identity, and autonomy.

Interface Zone (I0). Band I0 is the singular chokepoint where biological and synthetic signals convert. For invasive BCIs, this is the electrode-tissue boundary where electrical signals are transduced from ionic currents in neural tissue to electronic currents in silicon. For non-invasive devices, it is the sensor surface (e.g., scalp electrodes for EEG). The interface band is the narrowest point in the hourglass and the most critical for security: every neural signal that reaches the synthetic system, and every stimulation signal that reaches neural tissue, must pass through I0.

Synthetic Zone (S1–S3). The synthetic zone covers electronic and wireless systems from the analog front-end through digital processing to radio transmission. S1 handles analog signal conditioning and near-field electromagnetic effects. S2 encompasses digital processing, Bluetooth Low Energy, WiFi, and telemetry protocols. S3 covers radio frequency transmission, directed energy, and application-layer communication. Threats in this zone are closest to traditional IT security and are most amenable to conventional countermeasures.

3.5 Security Properties

The hourglass architecture provides three properties relevant to security analysis:

1. **Completeness:** Every component of a BCI system maps to exactly one band. There is no attack surface that falls outside the model.
2. **Chokepoint identification:** Band I0 identifies the natural monitoring point where all signals can be inspected during transit between neural and synthetic zones.
3. **Threat locality:** Each technique in the TARA taxonomy (Section 4) is annotated with the bands it affects, enabling defenders to understand the spatial extent of an attack across the architecture.

Figure 1 illustrates the 11-band model with relative band widths reflecting attack surface breadth at each layer.

QIF Hourglass Architecture

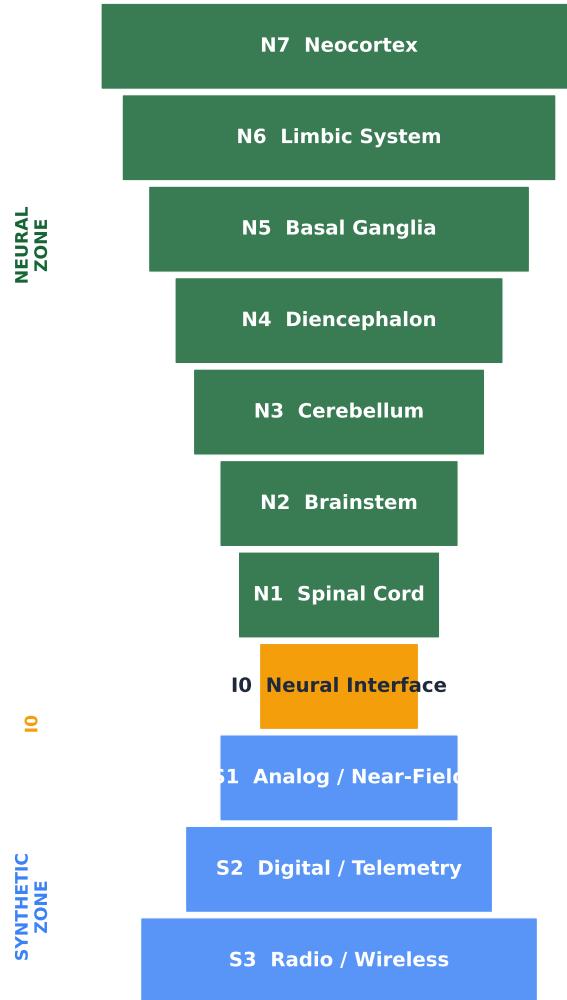


Figure 1: The QIF Hourglass Architecture. Band widths reflect relative attack surface breadth. The neural interface (I0) forms the narrow waist—the security chokepoint through which all signals must pass.

4 TARA: Threat Taxonomy

The Therapeutic Applications & Risk Assessment (TARA) registry catalogs BCI attack techniques with a structure inspired by MITRE ATT&CK [21], independently developed to cover neural, cognitive, and physical safety domains. Each technique is simultaneously an attack vector, an ethical risk, and—where applicable—a therapeutic application.

Why not ATT&CK directly? MITRE ATT&CK is designed for enterprise IT and mobile environments where adversary objectives center on data exfiltration, lateral movement, and persistence. BCI threats differ in three fundamental ways that make direct adoption inadequate: (1) the target is biological tissue, not an information system—attacks can cause seizures, tissue necrosis, or cognitive coercion, none of which map to ATT&CK’s impact categories; (2) the same technique is often simultaneously an attack and a therapy (e.g., deep brain stimulation is

both a Parkinson’s treatment and a neural injection vector), requiring a dual-use classification that ATT&CK has no mechanism to express; and (3) the attack lifecycle must span from radio-frequency transmission through silicon processing to neural tissue, crossing the bio-digital boundary that ATT&CK’s purely digital model does not address. TARA adopts ATT&CK’s proven structural methodology—techniques organized into tactics and domains—while building an independent taxonomy purpose-built for the BCI threat landscape.

4.1 Taxonomy Structure

The TARA registry organizes 102 techniques into 15 tactics across 8 operational domains. Technique identifiers follow the format **QIF-TXXX**, where the numeric suffix provides sequential ordering.

Domains. The eight domains span the full attack lifecycle:

1. **Neural (N)** — Direct interaction with neural tissue (scan, injection, manipulation)
2. **BCI System (B)** — System-level intrusion and evasion
3. **Protocol (P)** — Protocol disruption and communication attacks
4. **Data (D)** — Data harvesting and exfiltration
5. **Cognitive (C)** — Cognitive exploitation and imprinting
6. **Countermeasure (M)** — Surveillance and monitoring
7. **Evasion (E)** — Defense evasion and anti-detection
8. **Sensor (S)** — Consumer device side-channel attacks

Tactics. The 15 tactics follow a lifecycle structure analogous to ATT&CK, adapted for BCI operations: Neural Scan, BCI Intrusion, Neural Injection, Cognitive Imprinting, BCI Evasion, Data Harvesting, Neural Manipulation, Evasion/Rootkit Deployment, Monitoring/Surveillance, Protocol Disruption, Cognitive Exploitation, Signal Replay, Signal Harvesting, Signal Fingerprinting, and Signal Chaining.

4.2 Evidence Classification

Each technique carries an evidence status reflecting the maturity of its documentation:

Table 2: TARA evidence status classification with technique counts.

Status	Definition	Count
Confirmed	Documented in real-world use or peer-reviewed literature	19
Demonstrated	Proven in laboratory or controlled conditions	33
Emerging	Newly identified; limited but growing evidence	22
Theoretical	Plausible based on known physics and engineering	26
Plausible	Possible but with significant uncertainty	1
Speculative	Hypothetical; requires unproven capabilities	1

4.3 Severity Distribution

CVSS v4.0 base severity ratings for all 102 techniques show a distribution heavily weighted toward high and critical:

Table 3: TARA severity distribution (CVSS v4.0 base ratings).

Severity	Count	Percentage
Critical	29	28.4%
High	54	52.9%
Medium	16	15.7%
Low	3	2.9%

The dominance of high and critical ratings reflects the inherent severity of attacks against devices that interface directly with the nervous system. Even techniques with low exploitability can have severe consequences when the target is neural tissue.

4.4 Category Breakdown

The 102 techniques distribute across eight operational categories:

Table 4: TARA techniques by operational category.

ID	Category	Count
SE	Signal Eavesdropping	20
CI	Cognitive Integrity	18
EX	Data Exfiltration	17
DM	Data Manipulation	15
SI	Signal Injection	10
PE	Privilege Escalation	8
DS	Denial of Service	7
PS	Physical Safety	7

4.5 Dual-Use Mapping

A distinctive feature of TARA is the systematic dual-use mapping: every attack technique is assessed for therapeutic analogs. The same physical mechanisms that enable attacks—electromagnetic stimulation, signal decoding, neuromodulation—are the mechanisms underlying established therapies such as deep brain stimulation (DBS) [17], transcranial magnetic stimulation (TMS) [16], and neurofeedback.

Table 5: Dual-use classification of 102 TARA techniques.

Classification	Definition	Count	%
Confirmed	Published clinical use exists	52	51.0%
Probable	Under active clinical investigation	16	15.7%
Possible	Theoretical therapeutic mapping	9	8.8%
Silicon Only	No tissue analog; purely digital	25	24.5%

Of the 102 techniques, 77 (75.5%) have confirmed or probable therapeutic analogs. This finding underscores a fundamental challenge for BCI security: the same capabilities that must be defended against are often the capabilities that make BCIs therapeutically valuable.

4.6 Representative Techniques

Table 6 shows five representative techniques spanning different categories, severities, and dual-use classifications.

Table 6: Five representative TARA techniques.

ID	Sev.	Status	Dual-Use	Description
T0001	Critical	Confirmed	Confirmed	Cortical signal injection via rogue electrode stimulation; therapeutic analog: deep brain stimulation
T0015	High	Demonstrated	Confirmed	P300 side-channel extraction of private information; therapeutic analog: P300-based spelling interfaces
T0034	High	Emerging	Probable	Calibration data poisoning during BCI training sessions; therapeutic analog: adaptive neurofeedback
T0072	Medium	Confirmed	Confirmed	Ultrasonic side-channel via bone conduction microphone; therapeutic analog: ABR audiometry
T0090	High	Demonstrated	Confirmed	WiFi CSI body sensing for respiratory and gait inference; therapeutic analog: sleep apnea detection

Figure 2 shows the severity distribution across all techniques, and Figure 3 illustrates the dual-use breakdown.




Figure 2: TARA severity distribution across 102 techniques.

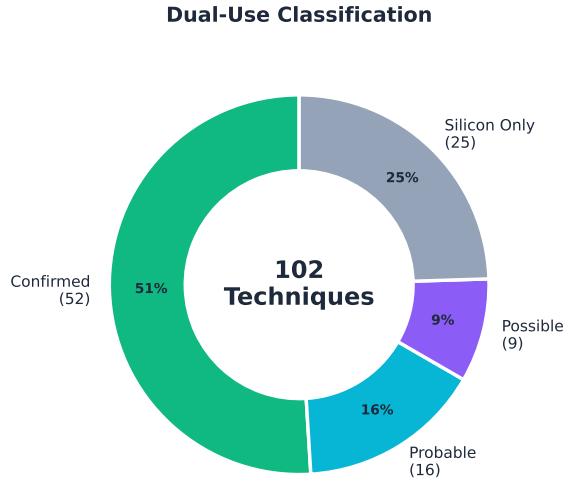


Figure 3: Dual-use classification: 75.5% of techniques have therapeutic analogs.

5 NISS: Neural Impact Scoring System

The Neural Impact Scoring System (NISS) is a proposed CVSS v4.0 extension designed to conform with FIRST.org’s official extension mechanism (User Guide §3.11) [10]. NISS adds five metrics that capture dimensions CVSS was never designed to express: biological tissue damage, cognitive integrity violations, consent boundary violations, damage reversibility, and neuroplastic consequences.

5.1 Gap Analysis: Why CVSS Alone Is Insufficient

We mapped all 102 TARA techniques to CVSS v4.0 base vectors and classified the results into three gap groups based on how much information CVSS alone fails to capture:

Table 7: CVSS v4.0 gap analysis across 102 TARA techniques.

Group	Gap Description	Count	Example
1	CVSS captures most impact; NISS adds nuance	12	Digital-only
2	CVSS captures exploitability but misses half	28	Mixed impact
3	CVSS fundamentally cannot express primary impact	58	Neural-dominant
Techniques needing NISS extension		98	96.1%

Group 3—where CVSS fundamentally cannot express the primary impact—contains the majority of techniques (56.9%). These are attacks where the most severe consequence is biological tissue damage, cognitive coercion, or irreversible neural harm—dimensions for which CVSS has no metric.

5.2 Extension Metrics

NISS defines five extension metrics, each with a graduated value set. The metrics are designed to be orthogonal to CVSS base metrics: they capture impact dimensions that exist only because the target system interfaces with biological neural tissue.

5.2.1 BI: Biological Impact

Direct harm to neural tissue, organs, or physiological function. This dimension has no equivalent in CVSS.

Value	Label	Score	Description
N	None	0.0	No tissue interaction or physical harm
L	Low	3.3	Temporary discomfort, minor sensory disruption, reversible tissue stress
H	High	6.7	Significant tissue damage, seizure induction, involuntary motor activation
C	Critical	10.0	Life-threatening or permanently disabling neural harm

5.2.2 CG: Cognitive Integrity

Impact on thought processes, perception, memory, identity, or decision-making. CVSS has no concept of thought privacy or cognitive autonomy.

Value	Label	Score	Description
N	None	0.0	No cognitive impact
L	Low	3.3	Decoded intent partially exposed, minor perceptual distortion
H	High	6.7	Full thought decoding, identity inference, or perception manipulation
C	Critical	10.0	Cognitive coercion, identity modification, or complete loss of cognitive autonomy

5.2.3 CV: Consent Violation

Degree of violation of informed consent or cognitive autonomy. Ordered by severity: covert (implicit) violations are worse than detectable (explicit) ones.

Value	Label	Score	Description
N	None	0.0	Operating within explicitly consented boundaries
P	Partial	3.3	Action exceeds scope but subject retains some awareness
E	Explicit	6.7	Direct violation of consent boundaries, but detectable
I	Implicit (covert)	10.0	Covert manipulation the patient cannot detect or refuse

5.2.4 RV: Reversibility

Whether the damage can be undone. IT assets can be restored from backup. Neural tissue cannot be rebooted.

Value	Label	Score	Description
F	Full	0.0	Effects fully reverse when attack stops
T	Temporary	3.3	Effects reverse over hours to days
P	Partial	6.7	Some effects permanent, some reversible
I	Irreversible	10.0	Permanent neural tissue destruction or cognitive change

5.2.5 NP: Neuroplasticity

Whether the attack exploits or induces neuroplastic changes—the brain’s ability to rewire itself. This has no digital equivalent.

Value	Label	Score	Description
N	None	0.0	No neuroplastic effect
T	Temporary	5.0	Short-term synaptic changes that decay within hours–days
S	Structural	10.0	Long-term or permanent neural pathway changes

5.3 PINS Flag

NISS introduces the Potential Impact to Neural Safety (PINS) flag—a binary indicator triggered when:

$$\text{PINS} = \begin{cases} \text{true} & \text{if BI} \geq \text{High} \vee \text{RV} = \text{Irreversible} \\ \text{false} & \text{otherwise} \end{cases} \quad (1)$$

A PINS flag mandates immediate safety review regardless of overall score. Across all 102 techniques, 31 are PINS-flagged (30.4%).

5.4 Scoring Formula

The NISS score is computed as the weighted mean of the five metric scores:

$$\text{NISS} = \frac{w_{\text{BI}} \cdot \text{BI} + w_{\text{CG}} \cdot \text{CG} + w_{\text{CV}} \cdot \text{CV} + w_{\text{RV}} \cdot \text{RV} + w_{\text{NP}} \cdot \text{NP}}{w_{\text{BI}} + w_{\text{CG}} + w_{\text{CV}} + w_{\text{RV}} + w_{\text{NP}}} \quad (2)$$

In the default profile, all weights are 1.0, yielding a simple arithmetic mean. NISS supports four context profiles with differential weights:

- **Clinical:** Emphasizes BI, RV, and NP (patient safety focus)
- **Research:** Emphasizes CG and CV (consent and cognition focus)
- **Consumer:** Balanced weights (general-purpose)
- **Military:** Emphasizes BI and CG (dual-use concern)

5.5 Vector Format

The NISS vector rides alongside the CVSS v4.0 base vector:

```
CVSS : 4.0 / AV : N / AC : L / AT : N / PR : N / UI : N /
      VC : H / VI : H / VA : H / SC : N / SI : N / SA : N
NISS : 1.0 / BI : H / CG : C / CV : I / RV : P / NP : S
```

This dual-vector architecture means security teams can triage using familiar CVSS scores while BCI-specific teams see the neural dimensions that determine whether a vulnerability is a software bug or a patient safety emergency.

5.6 NISS Severity Distribution

Across all 102 techniques, the NISS severity distribution differs markedly from CVSS severity:

Table 8: NISS severity distribution (all 102 techniques).

NISS Severity	Count	%
High	21	20.6%
Medium	29	28.4%
Low	51	50.0%
None	1	1.0%

The NISS distribution is more uniform than CVSS because NISS captures impact dimensions that CVSS flattens into high/critical. Techniques rated “high” by CVSS may distribute across low, medium, and high NISS scores depending on whether they involve biological tissue damage or are purely digital.

Figure 4 visualizes the gap between CVSS and NISS scoring.




Figure 4: CVSS v4.0 vs. NISS gap analysis: 96.1% of techniques require extension metrics CVSS cannot express.

6 Neural Impact Chain

The Neural Impact Chain (NIC) is a six-stage methodology for mapping security vulnerabilities to clinical psychiatric diagnoses. To our knowledge, this is the first systematic pipeline connecting cybersecurity severity scoring to DSM-5-TR diagnostic codes [1]. The NIC answers a question no prior framework has addressed: *if this attack succeeds, what psychiatric condition could it cause or worsen?*

6.1 Pipeline Architecture

The NIC traces each technique through six stages:

1. **Technique:** The TARA attack technique (e.g., cortical signal injection)
 2. **Hourglass Band:** Which band(s) the technique affects (e.g., N7 Neocortex, N6 Limbic)
 3. **Neural Structure:** The anatomical structure at risk (e.g., hippocampus, prefrontal cortex, amygdala)
 4. **Cognitive Function:** The function that structure supports (e.g., memory consolidation, executive control, emotional regulation)
 5. **NISS Score:** The neural impact score, particularly the BI, CG, and NP metrics that correlate with clinical outcomes
 6. **DSM-5-TR Code:** The ICD-10-CM diagnostic code(s) for the psychiatric condition most closely associated with disruption of that function
- Figure 5 illustrates this pipeline.

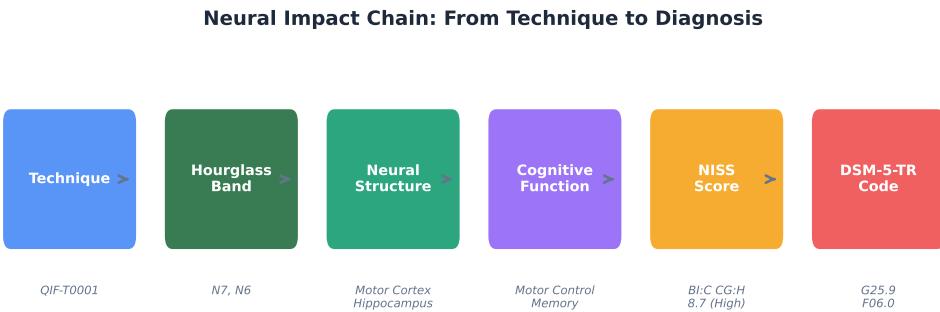


Figure 5: The Neural Impact Chain: six-stage pipeline from security technique to psychiatric diagnosis. Each arrow represents a mapping grounded in neuroanatomy, functional neuroscience, or clinical psychiatry.

6.2 NISS-to-DSM Bridge

The bridge between NISS metrics and DSM-5-TR diagnostic clusters is driven by which NISS metric dominates the technique’s profile:

- **BI-driven** → Motor/Neurocognitive cluster: biological impact implies tissue damage, leading to movement disorders or neurocognitive deficits
- **CG-driven** → Cognitive/Psychotic cluster: cognitive integrity violations imply perceptual or thought-process disruption
- **CV-driven** → Mood/Trauma cluster: Consent violations imply autonomy loss, mapping to trauma- and stressor-related disorders
- **NP/RV-driven** → Persistent/Personality cluster: neuroplasticity exploitation or irreversible damage implies lasting personality changes
- **No neural impact** → Non-diagnostic: Silicon-only techniques with no direct psychiatric mapping

6.3 Coverage Statistics

All 102 TARA techniques have been mapped through the NIC pipeline:

Table 9: Neural Impact Chain mapping results across 102 techniques.

Metric	Value
Techniques mapped	102 / 102 (100%)
Unique DSM-5-TR codes	15
Diagnostic clusters	5
Direct diagnostic risk	51 (50.0%)
Indirect diagnostic risk	9 (8.8%)
No diagnostic risk (silicon-only)	42 (41.2%)

6.4 Diagnostic Cluster Distribution

The five diagnostic clusters and their technique counts:

Table 10: DSM-5-TR diagnostic cluster distribution.

Cluster	Count	Representative DSM-5-TR Codes
Non-diagnostic	42	— (silicon-only techniques)
Mood/Trauma	21	F43.10 (PTSD), F32.9 (MDD), F44.9
Cognitive/Psychotic	16	F06.0 (psychosis), R41.3 (cognitive decline)
Motor/Neurocognitive	16	G25.9 (movement disorder), G31.84
Persistent/Personality	7	F07.0 (personality change)

The Mood/Trauma cluster is the largest diagnostic cluster (21 techniques), reflecting the prevalence of consent-violation and autonomy-disruption attacks in the BCI threat landscape. Techniques that covertly manipulate neural signals without the subject’s knowledge or consent map naturally to trauma- and stressor-related disorders.

6.5 Risk Classification

Each technique is classified by diagnostic risk:

- **Direct** (51 techniques, 50.0%): The attack mechanism can directly trigger or worsen the mapped psychiatric condition. Example: forced cortical stimulation causing seizures maps to epilepsy-related diagnostic codes.
- **Indirect** (9 techniques, 8.8%): The attack creates downstream conditions that may lead to the diagnosis. Example: sustained data exfiltration of private thoughts causing anxiety does not directly produce the anxiety disorder but creates the conditions for it.
- **None** (42 techniques, 41.2%): Silicon-only techniques with no direct neural interaction and thus no psychiatric diagnostic mapping.

6.6 Example Walkthrough

Consider technique **QIF-T0001: Cortical Signal Injection**—direct injection of adversarial signals via rogue electrode stimulation.

1. **Technique:** QIF-T0001 (cortical signal injection)
2. **Band:** N7 (Neocortex), N6 (Limbic System)
3. **Structure:** Primary motor cortex (M1), prefrontal cortex (PFC), hippocampus
4. **Function:** Motor control, executive function, memory consolidation
5. **NISS:** BI:C / CG:H / CV:I / RV:P / NP:S → Score: 8.7 (High), PINS flagged
6. **DSM-5-TR:** G25.9 (movement disorder NOS), F06.0 (psychotic disorder due to another medical condition), F07.0 (personality change due to another medical condition)
7. **Risk class:** Direct—the stimulation itself can trigger seizures, involuntary movement, and perception distortion

Figure 6 shows the distribution of techniques across diagnostic clusters.

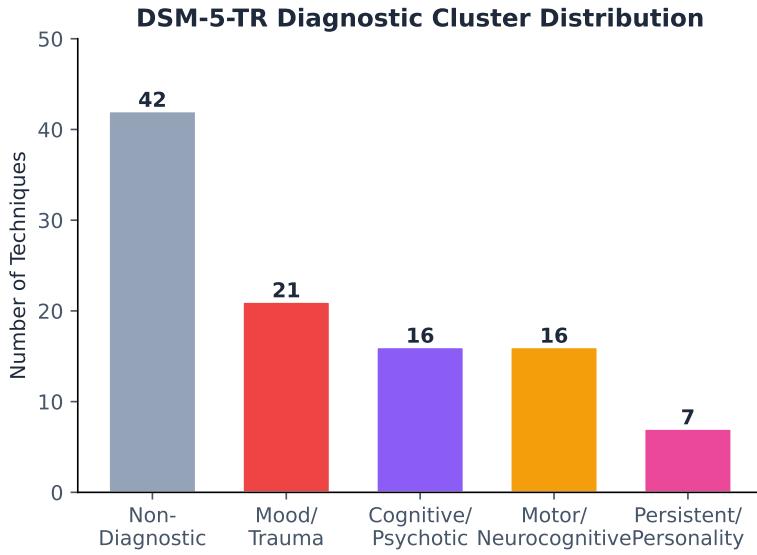


Figure 6: DSM-5-TR cluster distribution across 102 TARA techniques. 60 techniques (58.8%) have direct or indirect diagnostic risk.

7 Governance and Neuroethics

The QIF framework integrates neuroethics as a foundational design constraint rather than an afterthought. This section describes the consent tier system, alignment with international policy instruments, and regulatory mapping.

7.1 Consent Tiers

Each TARA technique is classified into one of four consent tiers based on the level of regulatory oversight required:

Table 11: Consent tier classification across 102 techniques.

Tier	Description	Requirement
Standard	Normal informed consent sufficient	Standard clinical consent
Enhanced	Additional safeguards required	Extended disclosure, monitoring
IRB	Institutional review board approval	Full IRB/ethics committee review
Prohibited	Not permissible under any consent	Technique must not be deployed

The distribution across tiers reflects the spectrum of BCI operations from routine monitoring (standard consent) to techniques that inherently violate cognitive autonomy (prohibited).

7.2 UNESCO Alignment

UNESCO’s 2025 Recommendation on the Ethics of Neurotechnology [30]—adopted by 194 Member States—is the first global normative framework for neurotechnology governance. It establishes three pillars: core values, ethical principles, and policy action areas.

The QIF framework addresses 15 of 17 UNESCO elements through technical implementation:

Table 12: QIF alignment with UNESCO Recommendation elements.

UNESCO Element	Type	Status	QIF Component
Human rights & dignity	Value	Implemented	Neurorights framework, consent tiers
Health & well-being	Value	Implemented	PINS flag, biological impact scoring
Diversity	Value	Implemented	Open-source, multi-stakeholder model
Sustainability	Value	Implemented	Apache 2.0 license, post-trial access
Professional integrity	Value	Implemented	Transparency audit trail
Proportionality	Principle	Implemented	Graduated severity scoring
Freedom of thought	Principle	Implemented	Cognitive liberty, consent states
Privacy	Principle	Implemented	Cognitive integrity metric (CG)
Protection of children	Principle	Implemented	Age-tiered consent framework
Consumer protection	Policy	Implemented	Default-deny architecture
Enhancement regulation	Policy	Partial	Technical infrastructure only
Workplace protections	Policy	Implemented	Mental privacy scoring
Behavioral influence	Policy	Implemented	Consent violation metric (CV)
Health and well-being	Policy	Implemented	NISS scoring, PINS flag
Oversight & governance	Policy	Implemented	Full regulatory mapping
Access & equity	Policy	Implemented	Open-source release

The two partially implemented elements—enhancement regulation and detailed implementation guidance for Member States—are explicitly outside the scope of a technical security framework and require policy collaboration.

7.3 Neurorights Framework Integration

The QIF framework implements the four neurorights proposed by Ienca and Andorno [12]:

1. **Cognitive Liberty:** Captured by the consent violation metric (CV). Any technique that operates without informed consent scores $CV \geq \text{Partial}$.
2. **Mental Privacy:** Captured by the cognitive integrity metric (CG). Techniques that decode intent, extract memories, or infer identity score $CG \geq \text{Low}$.
3. **Mental Integrity:** Captured by the biological impact metric (BI) and neuroplasticity metric (NP). Physical harm to neural tissue or induced structural changes violate mental integrity.
4. **Psychological Continuity:** Captured by the reversibility metric (RV). Irreversible changes to neural function threaten the continuity of personal identity.

7.4 Regulatory Mapping

The framework maps to existing and emerging regulatory instruments:

- **FDA Section 524B** [33]: Cybersecurity requirements for connected medical devices; NISS extends CVSS scoring as required by FDA premarket submissions
- **EU MDR 2017/745** [8]: Medical device risk management; TARA provides the threat taxonomy required for conformity assessment
- **ISO 14971** [14]: Risk management for medical devices; the NIC pipeline formalizes the harm pathway from technical failure to patient injury
- **HIPAA** [31]: Health data privacy; neural data classification extends protected health information categories
- **GDPR** [7]: Data protection; neural data as sensitive personal data under Article 9
- **Colorado Privacy Act** [29]: First US state to classify neural data as sensitive personal data (2024)

8 Case Studies

To demonstrate the practical difference between CVSS-only and CVSS+NISS scoring, we present five representative techniques scored with both systems. Each case illustrates dimensions that CVSS cannot capture.

Important caveat: Cases 1–4 are *threat model scenarios* derived from the TARA taxonomy. They represent plausible attack vectors based on known neuroscience and engineering principles, but have not been empirically executed against real BCI hardware. Case 5 is the only empirically confirmed vulnerability. This distinction is discussed further in Section 9.

8.1 Scoring Comparison

Table 13 presents five techniques spanning different categories, severity levels, and gap groups.

Table 13: CVSS v4.0 vs. NISS scoring for five representative techniques.

Technique	Description	CVSS	NISS
QIF-T0001	Cortical signal injection via rogue electrode	9.3 (Crit)	8.7 (High)
QIF-T0015	P300 side-channel private data extraction	7.7 (High)	4.7 (Med)
QIF-T0034	BCI calibration data poisoning	8.2 (High)	5.3 (Med)
QIF-T0050	Covert neural signal decoding for surveillance	7.1 (High)	6.0 (Med)
QIF-T0072	Ultrasonic bone-conduction side-channel	5.3 (Med)	2.7 (Low)

8.2 Case 1: Cortical Signal Injection (QIF-T0001)

CVSS v4.0 assessment. CVSS rates this technique as Critical (9.3) based on network attack vector, low complexity, no privileges required, and high impact on confidentiality, integrity, and availability. This accurately captures the exploitability and system impact.

What CVSS misses. The primary consequence of cortical signal injection is not system compromise—it is seizure induction, involuntary motor activation, and potential permanent tissue damage. The patient experiences a medical emergency, not a data breach. CVSS has no metric for biological harm, consent violation (the stimulation occurs without the patient’s knowledge), or irreversibility (neural tissue damage may be permanent).

NISS extension. NISS:1.0/BI:C/CG:H/CV:I/RV:P/NP:S — Score: 8.7 (High), PINS flagged. The NISS vector captures that this technique causes critical biological impact (BI:C), high cognitive integrity violation (CG:H), covert consent violation (CV:I), partially irreversible damage (RV:P), and structural neuroplastic changes (NP:S). The PINS flag triggers mandatory safety review.

NIC diagnostic mapping. Via the Neural Impact Chain: N7/N6 bands → motor cortex, hippocampus → motor control, memory → G25.9 (movement disorder), F06.0 (psychosis due to medical condition). Risk class: direct.

8.3 Case 2: P300 Side-Channel (QIF-T0015)

CVSS assessment. Rated High (7.7) for passive eavesdropping with high confidentiality impact.

What CVSS misses. The extracted data is not files or credentials—it is private cognitive responses. The P300 ERP component reveals whether the subject recognizes a stimulus, enabling extraction of PINs, personal preferences, and identity information [19]. CVSS treats this as a confidentiality breach; the actual impact is a cognitive integrity violation.

NISS extension. NISS:1.0/BI:N/CG:H/CV:E/RV:F/NP:N — Score: 4.7 (Medium). No biological impact, but high cognitive integrity violation and explicit consent violation. Fully reversible (passive attack). The NISS score is lower than CVSS because no physical harm occurs—but the CG:H flag alerts BCI-specific teams that private cognitive data is at risk.

8.4 Case 3: Calibration Poisoning (QIF-T0034)

CVSS assessment. Rated High (8.2) for integrity impact during the BCI training phase.

What CVSS misses. Poisoned calibration data causes the BCI to learn incorrect mappings between neural signals and intended actions. The patient’s device responds to wrong signals or fails to respond to correct ones. Over time, neuroplasticity causes the brain to adapt to the corrupted interface, creating lasting neural pathway changes even after the poisoning is discovered and corrected.

NISS extension. NISS:1.0/BI:L/CG:H/CV:I/RV:T/NP:T — Score: 5.3 (Medium). Low biological impact but high cognitive integrity violation (the device misinterprets intent), covert consent violation (the patient doesn’t know calibration was corrupted), and temporary neuroplastic changes.

8.5 Case 4: Covert Neural Surveillance (QIF-T0050)

CVSS assessment. Rated High (7.1) for sustained confidentiality impact.

What CVSS misses. Continuous covert decoding of neural signals constitutes ongoing mental privacy violation. Unlike data exfiltration from a server, the “data” being stolen is the patient’s thoughts, emotional states, and cognitive patterns. The consent violation is maximal: the patient cannot detect or refuse the surveillance.

NISS extension. NISS:1.0/BI:N/CG:C/CV:I/RV:F/NP:N — Score: 6.0 (Medium). No biological impact, but critical cognitive integrity violation (CG:C, full thought decoding) and covert consent violation (CV:I). The high CG and CV scores distinguish this from ordinary data exfiltration.

8.6 Case 5: Real-World Vulnerability Disclosure

The QIF framework has been applied to real vulnerability research. During systematic analysis of the BCI software ecosystem, we identified a multi-phase exploit chain in an open-source library used in clinical and research BCI pipelines.

The exploit chain demonstrates escalation from synthetic-zone vulnerabilities (S2/S3 bands) to potential neural-zone impact (N-band: corrupted data reaching clinical decision-making). CVSS scores the software vulnerabilities accurately; NISS captures the downstream risk to patients whose clinical care depends on the integrity of the data stream.

Responsible disclosure is in progress. Specific vulnerability details, including affected software and CWE identifiers, will be published after coordinated disclosure concludes.

9 Limitations and Future Work

We present these limitations transparently to guide future validation efforts and to prevent overstatement of the framework’s current maturity.

9.1 No Empirical Validation on Real BCI Devices

The QIF framework has not been validated against operational BCI hardware. The TARA taxonomy was developed through literature review, threat modeling, and systematic analysis rather than penetration testing of actual neural devices. While the framework has been applied to one real software vulnerability (Section 8), this covers only the synthetic zone. Validation against neural-zone and interface-zone threats requires access to implanted BCI patients and clinical environments—resources unavailable to independent researchers.

9.2 DSM-5-TR Mapping Not Clinically Validated

The Neural Impact Chain maps security techniques to psychiatric diagnoses based on known neuroanatomical pathways and functional neuroscience. However, these mappings have not been reviewed or validated by psychiatrists or clinical neuroscientists. The mappings represent our best assessment of which diagnostic codes correspond to disruption of specific neural functions, but clinical validation is essential before these mappings can inform clinical decision-making.

9.3 NISS Weights Not Calibrated

The NISS scoring formula (Equation 2) uses equal weights (1.0) for all five metrics in the default profile. The four context profiles (Clinical, Research, Consumer, Military) propose differential weights, but these have not been calibrated against empirical data, expert elicitation, or clinical outcomes. Weight calibration requires:

- Expert panel scoring of representative scenarios
- Sensitivity analysis across weight configurations
- Correlation with observed clinical outcomes (when available)

9.4 No Interrater Reliability Study

NISS scores in the TARA registry were assigned by a single analyst (the author). No interrater reliability study has been conducted to assess whether independent scorers would assign the same metric values. CVSS interrater reliability is a known challenge [9]; NISS, with its novel neural-specific metrics, likely faces greater variability. A formal interrater reliability study with domain experts from both cybersecurity and neuroscience is needed.

9.5 Taxonomy Completeness

The TARA registry contains 102 techniques as of version 1.4. The BCI threat landscape is evolving rapidly, and additional techniques will emerge as: (a) new BCI devices reach market, (b) consumer neurotechnology proliferates, and (c) adversarial AI techniques advance. The current registry should be treated as a foundation, not a complete enumeration.

Of the 102 techniques, 26 are classified as Theoretical and 1 as Speculative—these have not been empirically demonstrated. While they are grounded in known physics and engineering principles, their practical feasibility remains unvalidated.

9.6 Single-Author Bias

The framework was developed by a single independent researcher. While multi-model AI verification was used throughout development (Claude, Gemini, ChatGPT), the architectural

decisions, scoring assignments, and clinical mappings reflect a single perspective. Peer review and multi-disciplinary collaboration are essential for maturation.

9.7 AI Tool Disclosure

In accordance with arXiv policy on AI-assisted research, we disclose the following. Large language models (Claude, Gemini, ChatGPT) were used during the development of this framework for: literature review assistance, code generation for data analysis and visualization tools, editorial review, and cross-validation of technical claims. All framework architecture, threat taxonomy design, scoring methodology, clinical mapping decisions, and research conclusions were human-directed and human-verified. AI-generated outputs were treated as drafts subject to manual review. The author takes full responsibility for all content in this paper, irrespective of how it was generated. A complete, auditable transparency log documenting every AI contribution, human decision, and verification step is maintained at <https://github.com/qinnovates/qinnovate/blob/main/governance/TRANSPARENCY.md>.

An earlier version of this preprint (v1.0) contained citation errors introduced during AI-assisted bibliography construction, including three fabricated entries. These were corrected in v1.1 through a two-pass independent verification audit. All references have been verified against their source publications via DOI resolution, author publication pages, and database lookup. A full post-mortem is available in the project repository. This revision (v1.2) corrects an internal percentage inconsistency and adds the author responsibility statement above.

9.8 Future Work

1. **Reference implementation:** A software tool that automates NISS scoring and NIC mapping for new techniques
2. **Clinical validation:** Collaboration with psychiatrists to validate DSM-5-TR mappings
3. **Interrater reliability:** Formal study with cybersecurity and neuroscience domain experts
4. **FIRST.org registration:** Formal submission and registration of NISS as a CVSS v4.0 extension
5. **Empirical testing:** Penetration testing against BCI hardware in controlled environments
6. **Weight calibration:** Expert elicitation and sensitivity analysis for NISS context profile weights
7. **Conference paper:** Condensed version for submission to Graz BCI Conference 2026 and USENIX WOOT '26

10 Conclusion

Brain-computer interfaces are transitioning from laboratory prototypes to commercial medical devices. The security frameworks designed for information technology—while necessary—are insufficient for devices that read and write neural signals. A vulnerability in a BCI is not merely a software bug; it is a potential path to seizures, cognitive manipulation, privacy violation at the level of thought, and irreversible neural harm.

This paper presented the QIF framework: an integrated system comprising an 11-band hourglass architecture, a 102-technique threat taxonomy (TARA), a CVSS v4.0 extension for neural-specific scoring (NISS), and the Neural Impact Chain—a first-of-its-kind methodology for mapping security vulnerabilities to DSM-5-TR psychiatric diagnoses. Analysis of all 102 techniques reveals that 96.1% require scoring dimensions CVSS cannot express, 75.5% have therapeutic dual-use analogs, and 58.8% pose direct or indirect psychiatric diagnostic risk.

The framework has significant limitations—no empirical validation on BCI hardware, no clinical validation of DSM-5-TR mappings, and single-author scoring—which we have documented transparently. These are not reasons to delay publication; they are invitations to collaborate.

The BCI industry is moving faster than security standards. Neuralink, Synchron, and Blackrock Neurotech are implanting devices in patients today. The gap between what these devices can do and what security frameworks can assess grows wider each month.

The complete framework, threat registry, NISS specification, and scoring data are released as open source under the Apache 2.0 license. We invite collaboration from the neuroscience, neuroethics, cybersecurity, and clinical psychiatry communities. Formal registration of NISS with FIRST.org's CVSS Special Interest Group is planned as future work.

The question is no longer whether BCI security frameworks are needed. The question is whether they will be ready before the first patient is harmed.

References

- [1] American Psychiatric Association. *Diagnostic and Statistical Manual of Mental Disorders, Fifth Edition, Text Revision (DSM-5-TR)*. American Psychiatric Association Publishing, 2022. doi: 10.1176/appi.books.9780890425787.
- [2] Rebecca Bellan. OpenAI invests in Sam Altman’s brain computer interface startup Merge Labs. TechCrunch, January 2026. \$252M seed round at \$850M valuation; co-founded by Sam Altman, Alex Blania, and Caltech researchers.
- [3] Tamara Bonaci, Ryan Calo, and Howard J. Chizeck. App stores for the brain: Privacy & security in brain-computer interfaces. *IEEE Technology and Society Magazine*, 34(2):32–39, 2015. doi: 10.1109/MTS.2015.2425551.
- [4] Carmen Camara, Pedro Peris-Lopez, and Juan E. Tapiador. Security and privacy issues in implantable medical devices: A comprehensive survey. *Journal of Biomedical Informatics*, 55:272–289, 2015. doi: 10.1016/j.jbi.2015.04.007.
- [5] Stephen Deering. Watching the waist of the protocol hourglass. Keynote, IETF 51 Plenary, 2001. Originally presented at ICNP ’98. Foundational description of the Internet hourglass architecture.
- [6] Tamara Denning, Yoky Matsuoka, and Tadayoshi Kohno. Neurosecurity: Security and privacy for neural devices. *Neurosurgical Focus*, 27(1):E7, 2009. First use of the term “neurosecurity” in academic literature.
- [7] European Parliament and Council. General data protection regulation (GDPR) — regulation (EU) 2016/679, 2016. URL <https://gdpr-info.eu/>.
- [8] European Parliament and Council. Regulation (EU) 2017/745 on medical devices (MDR), 2017. URL <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32017R0745>.
- [9] FIRST.org. Common vulnerability scoring system v4.0 specification, 2023. URL <https://www.first.org/cvss/v4-0/cvss-v40-specification.pdf>.
- [10] FIRST.org. CVSS v4.0 user guide, 2023. URL <https://www.first.org/cvss/v4.0/user-guide>. Section 3.11: Extension Framework.
- [11] Daniel Halperin, Thomas S. Heydt-Benjamin, Benjamin Ransford, et al. Pacemakers and implantable cardiac defibrillators: Software radio attacks and zero-power defenses. In *Proceedings of the 2008 IEEE Symposium on Security and Privacy*, pages 129–142, 2008. doi: 10.1109/SP.2008.31.

- [12] Marcello Ienca and Roberto Andorno. Towards new human rights in the age of neuroscience and neurotechnology. *Life Sciences, Society and Policy*, 13(1):5, 2017. doi: 10.1186/s40504-017-0050-1.
- [13] Marcello Ienca and Pim Haselager. Hacking the brain: Brain-computer interfacing technology and the ethics of neurosecurity. *Ethics and Information Technology*, 18(2):117–129, 2016. doi: 10.1007/s10676-016-9398-9.
- [14] International Organization for Standardization. ISO 14971:2019 — medical devices — application of risk management to medical devices, 2019.
- [15] Eric R. Kandel, John D. Koester, Sarah H. Mack, and Steven A. Siegelbaum. *Principles of Neural Science*. McGraw-Hill, 6th edition, 2021.
- [16] Joachim K. Krauss, Nir Lipsman, Tipu Aziz, et al. Technology of deep brain stimulation: Current status and future directions. *Nature Reviews Neurology*, 17:75–87, 2021. doi: 10.1038/s41582-020-00426-z.
- [17] Andres M. Lozano, Nir Lipsman, Hagai Bergman, et al. Deep brain stimulation: Current challenges and future directions. *Nature Reviews Neurology*, 15:148–160, 2019. doi: 10.1038/s41582-018-0128-2.
- [18] Gabriel Lázaro-Muñoz, Michelle T. Pham, Katrina A. Muñoz, et al. Post-trial access in implanted neural device research. *Brain Stimulation*, 15(5):1029–1036, 2022. doi: 10.1016/j.brs.2022.07.051.
- [19] Ivan Martinovic, Doug Davies, Mario Frank, Daniele Perito, Tomas Ros, and Dawn Song. On the feasibility of side-channel attacks with brain-computer interfaces. In *Proceedings of the 21st USENIX Security Symposium*, pages 143–158. USENIX Association, 2012.
- [20] Lubin Meng, Xue Jiang, and Dongrui Wu. Adversarial robustness benchmark for EEG-based brain-computer interfaces. *Future Generation Computer Systems*, 2023. doi: 10.1016/j.future.2023.01.028.
- [21] MITRE Corporation. ATT&CK® framework, 2024. URL <https://attack.mitre.org/>.
- [22] Elon Musk and Neuralink. An integrated brain-machine interface platform with thousands of channels. *Journal of Medical Internet Research*, 21(10):e16194, 2019. doi: 10.2196/16194.
- [23] Katrina A. Muñoz, Kristin Kostick, Clarissa Sanchez, Lavina Kalwani, Laura Torgerson, Rebecca Hsu, Demetrio Sierra-Mercado, Jill O. Robinson, Simon Outram, Barbara A. Koenig, Stacey Pereira, Amy McGuire, Peter Zuk, and Gabriel Lázaro-Muñoz. Researcher perspectives on ethical considerations in adaptive deep brain stimulation trials. *Frontiers in Human Neuroscience*, 14:578695, 2020. doi: 10.3389/fnhum.2020.578695.
- [24] OECD. Recommendation on responsible innovation in neurotechnology, 2019. URL <https://legalinstruments.oecd.org/api/print?ids=658&Lang=en>.
- [25] Thomas J. Oxley, Peter E. Yoo, Gill S. Rind, et al. Motor neuroprosthesis implanted with neurointerventional surgery improves capacity for activities of daily living tasks in severe paralysis. *Journal of NeuroInterventional Surgery*, 13:102–108, 2021. doi: 10.1136/neurintsurg-2020-016862.
- [26] Republic of Chile. Constitutional amendment on neurorights (article 19, no. 1), 2021. First country to constitutionally protect neurorights.

- [27] Michael Rushanan, Aviel D. Rubin, Denis Foo Kune, and Colleen M. Swanson. SoK: Security and privacy in implantable medical devices and body area networks. In *Proceedings of the 2014 IEEE Symposium on Security and Privacy*, pages 524–539, 2014. doi: 10.1109/SP.2014.40.
- [28] Tyler Schroder, Renee Sirbu, Sohee Park, Jessica Morley, Sam Street, and Luciano Floridi. Cyber risks to next-gen brain-computer interfaces: Analysis and recommendations. *arXiv preprint arXiv:2508.12571*, 2025.
- [29] State of Colorado. Colorado privacy act — amendment on biological and neural data, 2024. First US state to classify neural data as sensitive personal data.
- [30] UNESCO. Recommendation on the ethics of neurotechnology. Adopted at the 43rd session of the General Conference, Samarkand, 2025. URL <https://www.unesco.org/en/ethics-neurotech/recommendation>.
- [31] U.S. Congress. Health insurance portability and accountability act (HIPAA), 1996.
- [32] U.S. Food and Drug Administration. Cybersecurity in medical devices: Quality system considerations and content of premarket submissions, 2023. URL <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/cybersecurity-medical-devices-quality-management-system-considerations-and-content-pre> Superseded Feb 2026 by updated guidance.
- [33] U.S. Food and Drug Administration. Section 524B of the FD&C act — ensuring cybersecurity of medical devices, 2023. Enacted via the Consolidated Appropriations Act, 2023.
- [34] Francis R. Willett, Erin M. Kunz, Chaofei Fan, et al. A high-performance speech neuroprosthesis. *Nature*, 620:1031–1036, 2023. doi: 10.1038/s41586-023-06377-x.
- [35] Rafael Yuste, Sara Goering, Blaise Agüera y Arcas, et al. Four ethical priorities for neurotechnologies and AI. *Nature*, 551(7679):159–163, 2017. doi: 10.1038/551159a.
- [36] Hubert Zimmermann. OSI reference model — the ISO model of architecture for open systems interconnection. *IEEE Transactions on Communications*, 28(4):425–432, 1980. doi: 10.1109/TCOM.1980.1094702.